DESIGN REVIEW BOARD

GACHŌ BADI (GOOSE BUDDY)

GDD Draft #1 · Reviewed 2026-07-23

> This document is not ready to drive production. The single change that matters most is resolving whether Gachō Badi is a finite, completable puzzle game or an endless social sandbox — because the economy math, the cost budget, the task-completion architecture, and the pitch's own marketing language all currently assume different, incompatible answers to that one question.

30 TOTAL FINDINGS

15 BLOCKING

2 UNRESOLVED DISAGREEMENTS

TOP 5 ISSUES

01

BLOCKING STRENGTHENED

The finite "complete all tasks" ending contradicts the endless-sandbox premise the pitch is sold on.

All six reviewers independently flagged this in Round 1. The Executive Summary sells "endless possibilities," but the document twice states the game has a hard stop once all tasks are completed — and the loop is stated to spawn more tasks as the island grows, so Systems Designer and Adversarial QA both question whether that stop is even reachable. Multiple reviewers upgraded their own severity after seeing how many colleagues converged on the same seam independently.

SYSTEMS DESIGNER | NARRATIVE CRITIC | PLAYER PSYCHOLOGIST | FEASIBILITY LEAD | ADVERSARIAL QA | BUSINESS ANALYST

02

BLOCKING | STRENGTHENED

No system in the architecture validates, detects, or gates task completion.

No listed agent confirms that an emergent, physics-driven chain of actions actually satisfied a task's goal state. Cross-examination showed this is the root cause of three separately-filed symptoms: an unsolvable-task risk, a credit-farming exploit, and a permanent softlock risk — all three reviewers reframed their findings as downstream of this one missing subsystem.

FEASIBILITY LEAD | SYSTEMS DESIGNER | ADVERSARIAL QA

03

BLOCKING | STRENGTHENED

The Director Agent's script-following mechanism is unexplained and structurally conflicts with the player-driven sandbox.

Feasibility Lead flags that grounding generated text into deterministic physical behavior is an open engineering problem treated as a one-line pipeline step. Adversarial QA independently flags that the player can interrupt any Director-driven scene at any moment, with no stated abort or replan behavior. Together these combine into a harder problem: real-time replanning against arbitrary player disruption, not a one-time translation.

FEASIBILITY LEAD | ADVERSARIAL QA

04

BLOCKING | STRENGTHENED

The document is unquantified where it matters and over-quantified where it can't be trusted.

Systems Designer shows the in-fiction credit economy has zero numbers. Business Analyst shows the one table that does exist — LLM token budgets — is explicitly disclaimed by the author as ChatGPT-generated and unvalidated, yet governs live recurring cost. The team quantified the variable it admits it doesn't understand while leaving the variable in its own control unquantified.

SYSTEMS DESIGNER | BUSINESS ANALYST | FEASIBILITY LEAD

05

BLOCKING | STRENGTHENED

No team size, timeline, budget, platform, audience, or monetization model is stated anywhere.

The only resourcing signal in the document is an aside that the author has student access to ChatGPT and is unfamiliar with costs. Feasibility Lead and Business Analyst independently flagged the identical missing section from the technical-buildability and market/greenlight angles respectively, and both converged on BLOCKING once they saw the overlap.

FEASIBILITY LEAD BUSINESS ANALYST

Notable runner-up, not in the top 5 by severity but converged upon by three reviewers: the pitch's stated emotional hook — resident "drama" and "relationships" — has no supporting mechanic (Narrative Critic), no tracking system (Adversarial QA), and no reward loop independent of the task economy (Player Psychologist). All three treated this as one missing layer once cross-examined, but none rated it above MAJOR.

THE DISAGREEMENT

What is the dominant player-facing failure mode — and it hinges on one undocumented rule.

PLAYER PSYCHOLOGIST

The unbounded, uncappable task backlog with no journal or wayfinding produces an overwhelm-and-quit scenario as the island grows.

ADVERSARIAL QA

The same repeatable physical actions (hose-drench, mustard-toss) are trivially farmable for infinite credit, producing boredom and exploit-driven trivialization instead.

Both reviewers flagged these as mutually exclusive predictions contingent on a single undocumented rule: whether a completed task is consumed or freely repeatable. Neither could resolve which applies without the designer answering that question first.

Does a content cap plus a post-completion sandbox mode actually fix Issue 01 — or just relocate the contradiction?

SYSTEMS DESIGNER & PLAYER PSYCHOLOGIST

Propose a mechanical fix: cap the content catalog so the completion state is reachable, and add a post-completion free-play mode so the sandbox promise isn't broken.

NARRATIVE CRITIC & BUSINESS ANALYST

A capped, converging loop still ends on a stated "harmony" nothing explains thematically. Capping the catalog also falsifies the "endless possibilities" marketing language, and an endless post-completion mode reopens the unbounded LLM-cost problem.

The board could not settle whether this is one fix or three separate fixes — systems, narrative, and cost — that happen to share a symptom.

FULL BOARD

SYSTEMS DESIGNER | NARRATIVE CRITIC | PLAYER PSYCHOLOGIST | FEASIBILITY LEAD | ADVERSARIAL QA | BUSINESS ANALYST

ROUND 1 FINDINGS

BLOCKING

The economy has no math

The progression currency is described with zero quantities — no credit yield per task, no unlock costs, no scaling. The Technical Strategy section quantifies LLM token budgets in detail but has no equivalent table for the actual game economy, so the loop cannot be evaluated, balanced, or implemented without at least placeholder numbers.

BLOCKING

The win condition and the expansion loop work against each other

The loop is additive and self-reinforcing — more residents/buildings spawn more tasks — but completion requires all tasks finished. The document never states whether the roles/building catalog is fixed or open-ended, so it's unclear whether "all tasks completed" ever stabilizes.

MAJOR

No stated mechanism guarantees a generated task is solvable

Tasks are open-ended, multi-step, physics/timing-dependent solutions. Nothing validates that a batch-generated task is achievable given the goose's verb set or current world state, and there's no fallback (hint, re-roll, expiry) if one isn't.

BLOCKING

Pacing and difficulty curve across a full playthrough are absent

The document names a tutorial and an ending but nothing in between describes how task complexity, resident count, or session length escalates, and tasks are stated to be non-sequential with no complexity tiering.

MAJOR

The loop has no cost, sink, or pushback

Systems are purely additive with no decay, expiring content, competing use for credits, or failure state. May be an intentional cozy-game choice, but it's never stated as a choice.

ROUND 2 — CROSS-EXAMINATION

CONFLICTS

Adversarial QA's credit-farming finding flips, rather than corroborates, my convergence concern — if tasks are freely repeatable, the economy collapses from the opposite direction (worthless credits, not unreachable completion). Both readings can't be true at once until the document states whether a task instance is consumed on completion.

Contested Player Psychologist's proposed fix (a post-completion sandbox mode) as insufficient on its own — an unbounded free-play phase with no economy math or sink just relocates the "everything is additive and eventually meaningless" problem to after the credits screen.

CONNECTIONS

My solvability gap and Feasibility Lead's missing completion-detection agent are two halves of the same failure: no guarantee a task is achievable at generation time, and no guarantee it's recognized as solved afterward — a full round-trip failure in the loop's most load-bearing verb.

Player Psychologist's unbounded task backlog (no journal, no wayfinding) sharpens my missing pacing curve into a concrete overwhelm scenario, not just an evaluability problem.

REVISIONS

Upgraded pacing/difficulty curve from MAJOR to BLOCKING — Player Psychologist's tutorial-to-main-game difficulty spike shows this produces plausible first-hour churn, not just a balancing headache.

Upgraded no-sink finding from MINOR to MAJOR — combined with QA's repeatable-task exploit and Feasibility's missing completion detection, the absence of any sink removes the last backstop against both farming and untunable economy.
